

Chapter 8

Deadlocks

Yunmin Go

School of CSEE



Agenda

- Introduction
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock

Deadlock State

- A set of threads is in a deadlocked state when every thread in the set is waiting for an event that can be caused only by another thread in the set
- The events with which we are mainly concerned here are resource acquisition and release




Deadlock in Multithreaded Application

■ Deadlock example

```
pthread_mutex_t first_mutex;  
pthread_mutex_t second_mutex;  
  
pthread_mutex_init(&first_mutex, NULL);  
pthread_mutex_init(&second_mutex, NULL);  
...
```

```
// thread_one runs in this function  
void *do_work_one(void *param) {  
    pthread_mutex_lock(&first_mutex);  
    pthread_mutex_lock(&second_mutex);  
    /**  
     * Do some work  
     */  
    pthread_mutex_unlock(&second_mutex);  
    pthread_mutex_unlock(&first_mutex);  
    pthread_exit(0);  
}
```

Resource-allocation graph




```
// thread_two runs in this function  
void *do_work_two(void *param) {  
    pthread_mutex_lock(&second_mutex);  
    pthread_mutex_lock(&first_mutex);  
    /**  
     * Do some work  
     */  
    pthread_mutex_unlock(&first_mutex);  
    pthread_mutex_unlock(&second_mutex);  
    pthread_exit(0);  
}
```

Deadlock Characterization

- Deadlock can arise if four conditions hold simultaneously
- **Mutual exclusion**: only one thread at a time can use a resource
- **Hold and wait**: a thread holding at least one resource is waiting to acquire additional resources held by other threads
- **No preemption**: a resource can be released only voluntarily by the thread holding it, after that thread has completed its task
- **Circular wait**: there exists a set $\{T_0, T_1, \dots, T_n\}$ of waiting threads such that T_0 is waiting for a resource that is held by T_1 , T_1 is waiting for a resource that is held by T_2 , ..., T_{n-1} is waiting for a resource that is held by T_n , and T_n is waiting for a resource that is held by T_0 .


Resource-Allocation Graph

- Deadlock can be described more precisely in terms of a directed graph called **resource-allocation graph**
- A set of vertices V and a set of edges E
- V : two types of vertices
 - $T = \{T_1, T_2, \dots, T_n\}$: a set of threads
 - $R = \{R_1, R_2, \dots, R_m\}$: a set of resource types
- E : two types of directed edges
 - $T_i \rightarrow R_j$: **request edge**
 - $R_j \rightarrow T_i$: **assignment edge**
 - If a request edge is fulfilled, it become assignment edge




Resource-Allocation Graph

- Example of resource-allocation graph
 - One instance of R_1
 - Two instances of R_2
 - One instance of R_3
 - Three instances of R_4
 - T_1 holds one instance of R_2 and is waiting for an instance of R_1
 - T_2 holds one instance of R_1 , one instance of R_2 , and is waiting for an instance of R_3
 - T_3 holds one instance of R_3




Resource-Allocation Graph

- If graph contains no cycles \rightarrow no deadlock
- If graph contains a cycle
 - If only one instance per resource type, then deadlock
 - If several instances per resource type, possibility of deadlock



<Resource-allocation graph with a deadlock>



<Resource-allocation graph with a cycle but no deadlock>

Methods for Handling Deadlocks

- Three ways to handle deadlock problem
 - Ignore the problem and pretend that deadlocks never occur in the system
 - Use a protocol to prevent or avoid deadlocks, ensuring that the system will never enter a deadlock state
 - Deadlock prevention
 - Deadlock avoidance
 - Allow the system to enter a deadlock state, detect it, and recover

Methods for Handling Deadlock

- Deadlock prevention

- A set of methods to ensure that at least one of **necessary conditions** cannot hold
- Constraint on how requests for resources

- Deadlock avoidance

- Keep the system in **safe state** in which deadlock cannot occur (using additional information)
 - Resources currently available or allocated to each thread
 - Additional information about future requests and release of each thread

Agenda

- Introduction
- **Deadlock Prevention**
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock

Deadlock Prevention

- Invalidate one of the four necessary conditions for deadlock:
- **Mutual exclusion**: not required for sharable resources (e.g., read-only files); must hold for non-sharable resources
 - Many resources are intrinsically non-sharable
- **Hold and wait**: must guarantee that whenever a thread requests a resource, it does not hold any other resources
 - Allocate all required resources before it begins execution or allow a thread to request resources only when it has none allocated to it
 - Low resource utilization and starvation possible

Deadlock Prevention

■ No Preemption

- If a thread that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released
- Preempted resources are added to the list of resources for which the thread is waiting
- Thread will be restarted only when it can regain its old resources, as well as the new ones that it is requesting

Deadlock Prevention

- **Circular Wait:** impose a total ordering of all resource types, and require that each thread requests resources in an increasing order of enumeration
 - Invalidating the circular wait condition is most common
 - Simply assign each resource (i.e. mutex locks) a unique number
 - Resources must be acquired in order
ex) $F(\text{first_mutex})=1$, $F(\text{second_mutex})=5$
- The thread can request an instance of resource R_j if and only if $F(R_j) > F(R_i)$
- code for thread_two could not be written as follows

```
// thread_one runs in this function
void *do_work_one(void *param) {
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);
    pthread_exit(0);
}
```

```
// thread_two runs in this function
void *do_work_two(void *param) {
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);
    pthread_exit(0);
}
```

Agenda

- Introduction
- Deadlock Prevention
- **Deadlock Avoidance**
- Deadlock Detection
- Recovery from Deadlock

Deadlock Avoidance


- Require additional information about how resources are to be requested.
 - ex) Each thread declares maximum # of resources of each type it may request
- Deadlock avoidance algorithm dynamically examines **resource-allocation state** to ensure a circular-wait condition can never exist
- **Resource-allocation state** is defined by
 - # of available resources
 - # of allocated resources
 - Maximum demands of threads

Safe State

- When a thread requests an available resource, system must decide if immediate allocation leaves the system in a **safe state**
- **Safe state**: there exists a **safe sequence** of all threads
- **Safe sequence**: a sequence $\langle T_1, T_2, \dots, T_n \rangle$ is safe if, for each T_i , the resources that T_i can still request can be satisfied by currently available resources plus the resources held by all T_j , with $j < i$
 - If the resources that T_i needs are not immediately available, then T_i can wait until all T_j have finished
 - When T_j is finished, T_i can obtain needed resources, execute, return allocated resources, and terminate
 - When T_i terminates, T_{i+1} can obtain its needed resources, and so on

Safe State and Deadlock

- If a system is in safe state → no deadlock
- If a system is in unsafe state → possibility of deadlock
- Avoidance → ensure that a system will never enter an unsafe state





Avoidance Algorithm

- Single instance of a resource type
 - Use a resource-allocation graph
- Multiple instances of a resource type
 - Use the Banker's Algorithm

Resource-Allocation Graph Algorithm

- **Claim edge** $T_i \rightarrow R_j$: thread T_i may request resource R_j at some time in the future; represented by a dashed line
 - Claim edge converts to request edge when a thread requests a resource
 - When a resource is released by a thread, assignment edge reconverts to a claim edge
 - Resources must be claimed a priori in the system
- The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph



Unsafe State!

Banker's Algorithm

- Multiple instances of resources
- Each thread must a priori claim maximum use
- When a thread requests a resource it may have to wait
- When a thread gets all its resources it must return them in a finite amount of time

Data Structures for the Banker's Algorithm

- Let n = number of threads, and m = number of resources types
 - **Available**: Vector of length m . If **Available** $[j]=k$, there are k instances of resource type R_j available
 - **Max**: $n \times m$ matrix. If **Max** $[i][j]=k$, then thread T_i may request at most k instances of resource type R_j
 - **Allocation**: $n \times m$ matrix. If **Allocation** $[i][j]=k$ then T_i is currently allocated k instances of R_j
 - **Need**: $n \times m$ matrix. If **Need** $[i][j]=k$, then T_i may need k more instances of R_j to complete its task
$$Need[i][j] = Max[i][j] - Allocation[i][j]$$

Safety Algorithm

1. Let ***Work*** and ***Finish*** be vectors of length m and n , respectively.
Initialize ***Work=Available*** and ***Finish[i]=false*** for $i=0,1,\dots,n-1$
2. Find an i such that both:
 - (a) ***Finish[i] = false***
 - (b) ***Need_i ≤ Work***If no such i exists, go to step 4
3. ***Work = Work + Allocation_i***
Finish[i] = true
Go to step 2
4. If ***Finish[i]==true*** for all i , then the system is in a safe state

Resource-Request Algorithm

$Request_i$ be the request vector for thread **T_i** .

If **$Request_i[j]=k$** then thread **T_i** wants **k** instances of resource type **R_j**

1. If **$Request_i \leq Need_i$** , go to step 2. Otherwise, raise error condition, since thread has exceeded its maximum claim.
2. If **$Request_i \leq Available$** , go to step 3. Otherwise **T_i** must wait, since resources are not available.
3. Pretend to allocate requested resources to **T_i** by modifying the state as follows:

$Available = Available - Request_i$

$Allocation_i = Allocation_i + Request_i$


$Need_i = Need_i - Request_i$

- If safe \rightarrow the resources are allocated to **T_i**
- If unsafe \rightarrow **T_i** must wait, and the old resource-allocation state is restored

Example: Banker's Algorithm

- 5 threads T_0 through T_4
- 3 resource types: A (10 inst.), B (5 inst.), and C (7 inst.)
- Snapshot at time t_0 :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>	<u>Need</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
T_0	0 1 0	7 5 3	3 3 2	7 4 3
T_1	2 0 0	3 2 2		1 2 2
T_2	3 0 2	9 0 2		6 0 0
T_3	2 1 1	2 2 2		0 1 1
T_4	0 0 2	4 3 3		4 3 1



- The system is in a safe state since the sequence $\langle T_1, T_3, T_4, T_2, T_0 \rangle$ satisfies safety criteria

Example: Banker's Algorithm

- Suppose T_1 requests (1, 0, 2)
 - Check that **Request** \leq **Available** (that is, $(1,0,2) \leq (3,3,2) \rightarrow \text{True}$)

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
T_0	0 1 0	7 4 3	2 3 0
T_1	3 0 2	0 2 0	
T_2	3 0 2	6 0 0	
T_3	2 1 1	0 1 1	
T_4	0 0 2	4 3 1	

- Executing safety algorithm shows that sequence $\langle T_1, T_3, T_4, T_0, T_2 \rangle$ satisfies safety requirement

Can request for (3,3,0) by T_4 be granted?

Can request for (0,2,0) by T_0 be granted?

Agenda


- Introduction
- Deadlock Prevention
- Deadlock Avoidance
- **Deadlock Detection**
- Recovery from Deadlock

Deadlock Detection


- Allow system to enter deadlock state
- Detection algorithm
 - Single instance of each resource type
 - Multiple instances of a resource type (skip, Additional slides)
- Recovery scheme

Single Instance of Each Resource Type

- **Wait-for graph:** removing resource nodes from resource-allocation graph and collapsing the appropriate edges




Resource-Allocation Graph



Corresponding wait-for graph

Single Instance of Each Resource Type

- A deadlock exists if and only if the wait-for graph contains a cycle.
- To detect deadlock, system should
 - Maintain wait-for graph
 - Invoke an algorithm to detect a cycle - $O(n^2)$



Wait-for graph

Detection-Algorithm Usage

- When should we invoke the detection algorithm?
 - How often is a deadlock likely to occur?
 - How many threads will be affected by deadlock when it happens?
- We may invoke detection algorithm whenever a request for allocation cannot be granted immediately
 - Deadlocks occur only when some threads makes a request that cannot be granted immediately
 - We can find the thread which caused deadlock.

Agenda

- Introduction
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock

Recovery from Deadlock

- Process termination
- Resource preemption

Process Termination

- Abort all deadlocked processes
 - Too expensive
- Abort one process at a time until the deadlock is eliminated
 - Order of priority
 - Time from start / time to completion
 - Resources the process has used / needs to complete
 - How many processes will need to be terminated?
 - Is the process interactive or batch?

Resource Preemption

- Selecting a victim
 - Minimizing cost (# of resources a process has, amount of time consumed so far, ...)
- Rollback
 - The selected process should return to some safe state and restart it.
- Starvation
 - How can we guarantee that resources will not always be preempted from the same process?

ADDITIONAL SLIDES

Several Instances of a Resource Type

- **Available:** A vector of length m indicates the number of available resources of each type
- **Allocation:** An $n \times m$ matrix defines the number of resources of each type currently allocated to each thread
- **Request:** An $n \times m$ matrix indicates the current request of each thread. If **Request** $[i][j] = k$, then thread T_i is requesting k more instances of resource type R_j .

Detection Algorithm

1. Let ***Work*** and ***Finish*** be vectors of length ***m*** and ***n***, respectively

Initialize:

(a) ***Work*** = ***Available***

(b) For $i = 1, 2, \dots, n$, if ***Allocation_i*** $\neq 0$, then
Finish[i] = ***false***; otherwise, ***Finish***[i] = ***true***

2. Find an index i such that both:

(a) ***Finish***[i] == ***false***

(b) ***Request_i*** \leq ***Work***

If no such i exists, go to step 4

Detection Algorithm

3. **$Work = Work + Allocation_i$** ;

$Finish[i] = true$

go to step 2

4. If **$Finish[i] == false$** , for some i , $1 \leq i \leq n$, then the system is in deadlock state. Moreover, if **$Finish[i] == false$** , then T_i is deadlocked

Algorithm requires an order of $O(m \times n^2)$ operations to detect whether the system is in deadlocked state

Example: Detection Algorithm

- Five threads T_0 through T_4
- Three resource types: A (7 inst.), B (2 inst.), and C (6 inst.)
- Snapshot at time t_0 :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
T_0	0 1 0	0 0 0	0 0 0
T_1	2 0 0	2 0 2	
T_2	3 0 3	0 0 0	
T_3	2 1 1	1 0 0	
T_4	0 0 2	0 0 2	

- Sequence $\langle T_0, T_2, T_3, T_1, T_4 \rangle$ will result in ***Finish[i] = true*** for all i

Example: Detection Algorithm

- T_2 requests an additional instance of type **C**

	<u>Request</u>			
	A	B	C	
T_0	0	0	0	
T_1	2	0	2	
T_2	0	0	1	
T_3	1	0	0	
T_4	0	0	2	

- State of system?
 - Can reclaim resources held by thread T_0 , but insufficient resources to fulfill other processes; requests
 - Deadlock exists, consisting of threads T_1 , T_2 , T_3 , and T_4